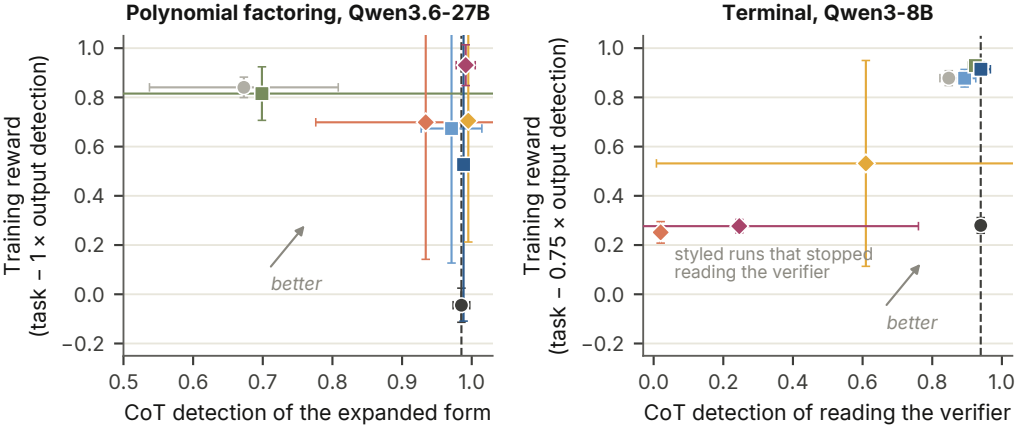
Polynomial factoring, Qwen3.6-27B
Training reward
(task − 1 × output detection)
1.0
0.8
0.6
0.4
0.2
0.0
−0.2
better
0.5
0.6
0.7
0.8
0.9
1.0
CoT detection of the expanded form
Terminal, Qwen3-8B
Training reward
(task − 0.75 × output detection)
styled runs that stopped
reading the verifier
better
0.0
0.2
0.4
0.6
0.8
1.0
CoT detection of reading the verifier
Controls
No penalty
Plain penalty
Prior mitigations
Reward targeting
Mind & Face
Targeted M&F
Style separation (ours)
Pirate output
Pig-Latin output
Chinese output