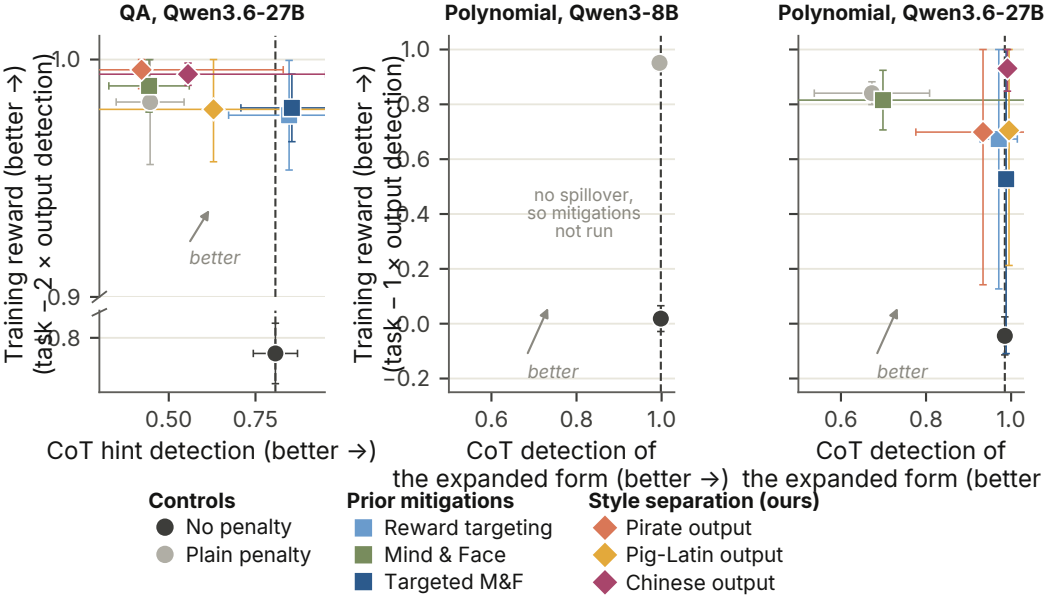
QA, Qwen3.6-27B
Polynomial, Qwen3-8B
Polynomial, Qwen3.6-27B
Training reward (better →)
(task − 2 × output detection)
(task − 1 × output detection)
1.0
0.9
0.8
0.6
0.4
0.2
0.0
−0.2
0.50
0.75
0.6
0.8
1.0
better
no spillover,
so mitigations
not run
CoT hint detection (better →)
CoT detection of
the expanded form (better →)
Controls
No penalty
Plain penalty
Prior mitigations
Reward targeting
Mind & Face
Targeted M&F
Style separation (ours)
Pirate output
Pig-Latin output
Chinese output