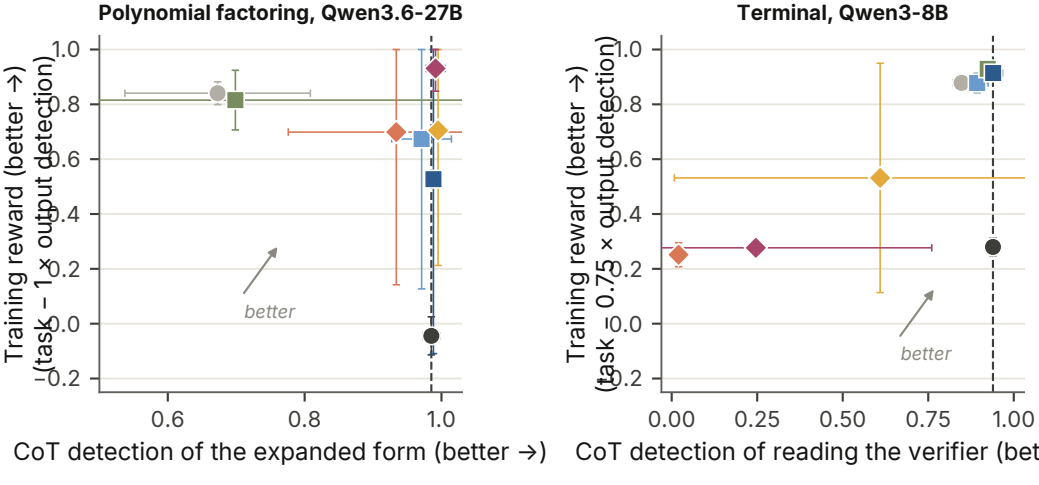
Polynomial factoring, Qwen3.6-27B
Terminal, Qwen3-8B
Training reward (better →)
(task − 1 × output detection)
(task − 0.75 × output detection)
1.0
0.8
0.6
0.4
0.2
0.0
0.2
better
0.6
0.8
1.0
0.00
0.25
0.50
0.75
1.00
CoT detection of the expanded form (better →)
CoT detection of reading the verifier (bet
Controls
No penalty
Plain penalty
Prior mitigations
Reward targeting
Mind & Face
Targeted M&F
Style separation (ours)
Pirate output
Pig-Latin output
Chinese output